

Building a Machine Learning Platform at Quora

The Quora Answer To “Build vs Buy” For ML Platforms

Nikhil Garg

@nikhilgarg28

A bit about me...

- At Quora since 2012
- Currently leading two ML engineering teams:
 - Content Quality
 - ML Platform



@nikhilgarg28

Quora

The best answer to any question



To Grow And Share World's Knowledge

What good things does the TPP hold for underemployed American citizens?

Will this trade agreement raise the U.S. workforce participation rate?

 Answer

Request ▼

Follow 6 Comment Share Downvote

...

9 Answers




Barack Obama, President of the United States

Written Thu · Featured in Inc · Upvoted by Marc Bodnick, [Harvard Gov major](#), [Stanford PoliSci PhD student](#), [Nadia Singer](#), and 27 others you follow

The underemployed – folks who work part-time and want to work full-time – have a lot to gain from the Trans-Pacific Partnership, or TPP. I'm not the only one who believes this – a number of independent, credible analyses of TPP have shown it will grow exports, grow our economy, and raise incomes.

Underemployment and unemployment are driven by an economy that's not producing at its full potential – factories that aren't humming at capacity and workers who aren't being put to their highest and best use. It makes sense – fewer goods. services. and exports

What are the best Caltrain hacks?



Follow

72

Comment

1

Share

Downvote



5 Answers



Matt Laroche, Ex-Caltrain Regular.

Updated Feb 25 · Upvoted by Jeremy Lipps, I am the Social Media Officer for Caltrain., Nadia Singer, and 23 others you follow



Tickets/Passes

- Load your Caltrain monthly (and other passes/credit) at Walgreen's (or another Clipper retail location: <https://clippercard.com/ClipperW...>). It avoids the 3 day wait when you buy a pass online. This also means you can use your monthly without doing the zone based tagging in and out. (There's a few edge cases that fail with online purchase of monthly passes)
- Monthly passes are valid for all zones on weekends and holidays!
- 2+ zone monthly passes are treated, basically, as monthly passes for VTA and

Over 100 million monthly uniques

Millions of questions & answers

In hundreds of thousands of topics

Supported by 80 engineers


What Slows Down ML Innovation?

Curse Of Complexity

- Pipeline jungles
- Lots of glue code to get data in/out of general purpose packages.
- Strong coupling between business logic, data, ML algorithms and configuration.




- Online vs offline
- Production vs experimentation
- C++ vs Python
- Engineering vs research
- ...even more glue code and pipeline jungles.



Getting New Applications Off The Ground

- Hard to reuse existing features, data, algorithms, tooling etc.
- Too costly to even get off the ground.



Machine Learning: The High-Interest Credit Card of Technical Debt

**D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov,
Todd Phillips, Dietmar Ebner, Vinay Chaudhary, Michael Young**
{dsculley, gholt, dgg, edavydov}@google.com
{toddphillips, ebner, vchaudhary, mwyoung}@google.com
Google, Inc

**One ring to bring them all and in
the darkness bind them!**

Machine Learning Platform

Collection of systems to sustainably increase the
business impact of ML at scale.

ML Platform: Build or Buy?

The Quora Answer: Build

For Seven Reasons

Reason # 7

Just Can't Buy Everything!

Degree Of Integration & Delegation



- No matter how powerful the platform is, still need to maintain some form of integration
- This thin integration layer then becomes the platform.
- Real questions --
 - How much does this in-house layer delegate?
 - How much control does it have over delegation?




Reason # 6

Fast Scalable Production Systems

End-To-End Online Production Systems




- External platforms at best can deploy “predictive models”, as services, not end-to-end online systems
- Gains come from optimizing the whole pipeline, not just algorithms.
- Latency: tens of milliseconds. Managing sharding, batching, data locality, caching, streaming, stragglers, graceful degradation...
- Real world systems -- boosts, diversity constraints, holes in data, skipping stages, hard filters... sounds familiar?



Reason # 5

Blurry Line Between Experimentation & Production

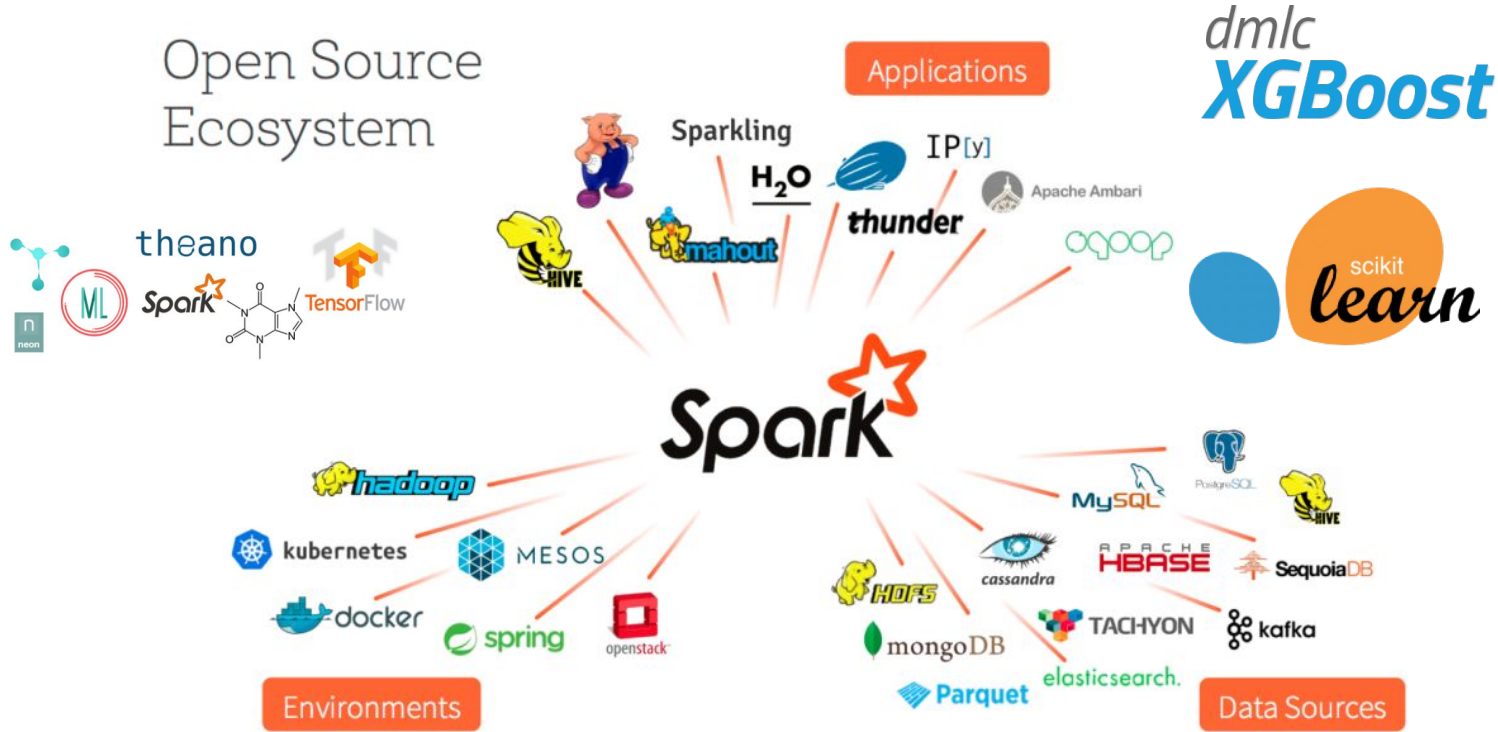
- We want the same code/systems/tools to work for both experimentation & production.
- But we need to carefully “control” the production code to keep it be fast.
- So need to “control” offline experimentation systems too.



Reason # 4

Openly Using Open Source


Open Source Ecosystem



Production ML Algorithms At Quora



- Logistic Regression
- Elastic Nets
- Random Forests
- Gradient Boosted Decision Trees
- Matrix Factorization
- (Deep) Neural Networks
- LambdaMart
- Clustering
- Random walk based methods
- Word Embeddings
- LDA
- ...



- Open source is great -- lots of great technologies!
- Commercial ML platforms are also open sourcing stuff.
- Learning and cherry-picking favorite parts from ANY open source systems.
- May write our own algorithms too (e.g QMF)
- Building own platform = controlling the delegation, not lack of delegation



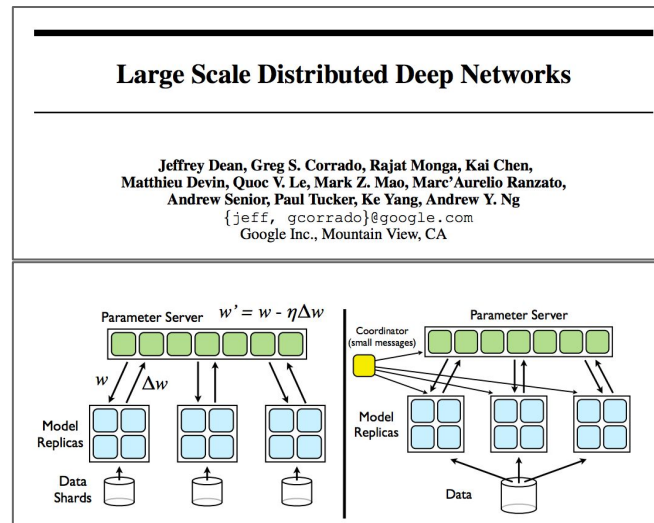
QMF - a matrix factorization library

build passing license Apache 2

Reason # 3

**Commercial Platforms' Offerings
Are Not Super Valuable To Us**

- Main offerings of external platforms are:
 - Lower operational overhead of running machines
 - Out-of-box distributed training.
- Operational overhead
 - Gets amortized over time
 - Shared with non-ML infrastructure.
- Can often train most models in a single multi-core machine.




Reason # 2

Blurry Line Between ML & Product Dev


Blurry Line Between ML/Non-ML Product



- Answer ranking
- Feed ranking
- Search ranking
- User recommendations
- Topic recommendations
- Duplicate questions
- Email Digest
- Request Answers
- Trending now
- Topic expertise prediction
- Spam, abuse detection
-




Discover new people




Shubham Bansal
DTU Mechanical Cla...
Answers about Cricket
(sport) and 7 more

Follow | 6.7k




Discover new topics



Gifted Education
Followed by Kat Tanaka
Okopnik

Follow | 10.6k



Top Stories from Your Feed
Your Quora Digest


Related Questions

Are there situations where L2 regularization is better than L1 regularization and vice versa?

What are the different ways of thinking about regularization?

Trending Now


- Indian Government Opposes Triple Talaq and Polygamy in SC
- Major Internet DDoS Attack

 Top Questions For You

Question asked · Leadership · 2h

Blurry Line Between ML/Non-ML Data

Quora



Billions of relationships and words

How would a Clinton administration help break down the gridlock in Washington?

[Answer](#) [Request](#) Follow 14 Comment 1 Share 10 Downvote

22 Answers



Hillary Clinton, Senator, Secretary of State, 2016 presidential candidate

I've been in and out of Washington for a long time, in a lot of different roles, and I know how much effort it takes to find common ground to get things done. There's nothing sexy about it. It's about getting up every day, building relationships—even with the people you don't agree with—and finding whatever sliver of common ground you can. From day one, that's exactly what I will do.

There's really no shortcut. You just have to keep pushing forward and keep listening and negotiating day after day. And no matter what, you have to keep reaching out. I've seen this work. It's how we were able to get the Children's Health Insurance Program passed when I was first lady. As a senator, I built alliances with people who were very much political adversaries to expand health care access for members of the National Guard and reservists. And as secretary of state, I rounded up Republicans to pass the New START Treaty. It takes a lot of effort, but if you're persistent, you can sit down across the table and across the aisle and find ways to get things done.

Upvoted 2.3k Downvote Comments 114+



Add a comment...

Comment

Ed Caruthers 110 votes Show


And we hope to give you a Democratic majority, at least in the Senate.

Blurry Line Between ML/Non-ML Codebase



- Integration with other utility libraries/services
e.g A/B testing, debug tools, monitoring, alerting, data transfer, ...
- Empowering all product engineers to do ML.

Variant	
True	Control
1,358,942	1,813,943
% CHANGE: +1.9%	
95% CI: (+0.8%, +3.0%)	




Reason # 1

ML As Quora's Core Competency

ML: Critical For Our Strategic Focus



- ML gives us a strategic competitive advantage.
- Want to control and develop deep expertise in the whole stack.
- Quora has a long term focus -- investment in platform more than pays off in the long term.
- Single most important reason to build ML Platform!



Summary

- Anyone doing non-trivial ML needs an ML platform to sustain innovation at scale.
- Build vs buy decision is not all-or-nothing.
- Surface area and importance of ML are deciding factors in the build vs buy decision.

Thank You!

YES, WE ARE HIRING :)

Nikhil Garg
@nikhilgarg28